

Problem: Find the shortest path between two nodes in a graph.

Constraints

- Is this a directional graph?
 - Yes
- Could the graph have cycles?
 - Yes
 - Note: If the answer were no, this would be a DAG.
 - DAGs can be solved with a [topological sort](#)
- Are the edges weighted?
 - Yes
 - Note: If the edges were not weighted, we could do a BFS
- Are the edges all non-negative numbers?
 - Yes
 - Note: Graphs with negative edges can be done with Bellman-Ford
 - Graphs with negative cost cycles do not have a defined shortest path
- Do we have to check for non-negative edges?
 - No
- Can we assume this is a connected graph?
 - Yes
- Can we assume the inputs are valid?
 - No
- Can we assume we already have a graph class?
 - Yes
- Can we assume we already have a priority queue class?
 - Yes
- Can we assume this fits memory?
 - Yes

Test Cases


The constraints state we don't have to check for negative edges, so we test with the general case.

- `graph.add_edge('a', 'b', weight=5)`
- `graph.add_edge('a', 'c', weight=3)`
- `graph.add_edge('a', 'e', weight=2)`
- `graph.add_edge('b', 'd', weight=2)`
- `graph.add_edge('c', 'b', weight=1)`
- `graph.add_edge('c', 'd', weight=1)`
- `graph.add_edge('d', 'a', weight=1)`
- `graph.add_edge('d', 'g', weight=2)`
- `graph.add_edge('d', 'h', weight=1)`

- `graph.add_edge('e', 'a', weight=1)`
- `graph.add_edge('e', 'h', weight=4)`
- `graph.add_edge('e', 'i', weight=7)`
- `graph.add_edge('f', 'b', weight=3)`
- `graph.add_edge('f', 'g', weight=1)`
- `graph.add_edge('g', 'c', weight=3)`
- `graph.add_edge('g', 'i', weight=2)`
- `graph.add_edge('h', 'c', weight=2)`
- `graph.add_edge('h', 'f', weight=2)`
- `graph.add_edge('h', 'g', weight=2)`
- `shortest_path = ShortestPath(graph)`
- `result = shortest_path.find_shortest_path('a', 'i')`
- `self.assertEqual(result, ['a', 'c', 'd', 'g', 'i'])`
- `self.assertEqual(shortest_path.path_weight['i'], 8)`

Algorithm

Wikipedia's animation:



Initialize the following:

- `previous = {}` # Key: node key, val: prev node key, shortest path
 - Set each node's previous node key to None
- `path_weight = {}` # Key: node key, val: weight, shortest path
 - Set each node's shortest path weight to infinity
- `remaining = PriorityQueue()` # Queue of node key, path weight
 - Add each node's shortest path weight to the priority queue
- Set the start node's `path_weight` to 0 and update the value in `remaining`
- Loop while `remaining` still has items
 - Extract the min node (node with minimum path weight) from `remaining`
 - Loop through each adjacent node in the min node
 - Calculate the new weight:

- Adjacent node's edge weight + the min node's path_weight
 - If the newly calculated path is less than the adjacent node's current path_weight:
 - Set the node's previous node key leading to the shortest path
 - Update the adjacent node's shortest path and update the value in the priority queue
- Walk backwards to determine the shortest path:
 - Start at the end node, walk the previous dict to get to the start node
- Reverse the list and return it

Complexity for array-based priority queue:

- Time: $O(v^2)$, where v is the number of vertices
- Space: $O(v^2)$

This might be better than the min-heap-based variant if the graph has a lot of edges.

$O(v^2)$ is better than $O((v + v^2) \log v)$.

Complexity for min-heap-based priority queue:

- Time: $O((v + e) \log v)$, where v is the number of vertices, e is the number of edges
- Space: $O((v + e) \log v)$

This might be better than the array-based variant if the graph is sparse.

-